


Operating Systems

Sistemas Operativos

CPU Scheduling

University of Minho
2024-2025




What will we learn?

CPU Management

- Although a computer has a limited number of CPUs, users get the impression that several programs (a lot more than the available CPUs) are running simultaneously
 - ▶ How does the OS provide this illusion?
 - ▶ How can users use the OS APIs to start and stop programs?
 - ▶ How does the OS choose what programs will be running, be switched, ...?
- Let us move to the **policies**!

Scheduler

- Number of processes ready to run is significantly higher than the number of available CPU cores
- The PCBs of *Ready* processes (check *Processes slides!*) are kept in a **Ready queue**
- The **OS scheduler** is responsible for choosing the next process to execute
 - ▶ By now, you should understand the mechanisms used by the OS to run another process
 - ▶ But what about the **policies**? What process should be scheduled next?



Workload Assumptions

Let's start simple...

- **Workload:** set of running processes
(Processes are also referred as **jobs** in the literature)
- Let's define some assumptions to reason about scheduling policies
 1. All processes run for the **same amount of time**¹, and the **time is known**
 2. All processes **arrive** in the system **at the same time**
 3. Once started, each process **runs until completion**
 4. Processes **only use the CPU** (i.e., no I/O is performed)
- Most of these assumptions are unrealistic...
...we will refine these as we move along

¹The amount of time that a process is using the CPU is also referred to as the process's **CPU Burst**

Metrics

How can we tell if algorithms are good?

- Let's define a first metric to compare scheduling algorithms

- ▶ **Turnaround time:** the time the process takes to complete after arriving in the system


$$T_{turnaround} = T_{completion} - T_{arrival}$$

- The turnaround time is a **performance** metric.
- There are other interesting non-performance metrics, for example
 - ▶ **CPU utilization**, as the OS wants to keep the CPU busy running useful instructions!
 - ▶ **Fairness**, i.e., if all processes get a chance to execute
 - ▶ ...

FCFS (aka FIFO)

First come, First Served (aka First in, First out)

- **Selection criteria:** Schedule the process that arrived first (earlier)
- **Example:** Processes P_1 , P_2 and P_3 , each taking 10 seconds to execute
 - ▶ **Order of arrival:** $P_1 \rightarrow P_2 \rightarrow P_3$
 - ▶ (we assume that all arrive at instant $t = 0$, with a very small delay among them)
- ▶ **Average turnaround time:** $\frac{10 + 20 + 30}{3} = 20 \text{ seconds}$



FCFS (aka FIFO)

Processes with distinct runtimes


- Lets relax Assumption 1: *all processes run for the same amount of time*
- **Example:** P_1 runs for 100 seconds, while P_2 and P_3 run for 10 seconds
- **Questions** (note that our goal is *minimizing the average turnaround time*)
 - ▶ Can you think of a workload where FCFS performs well?
 - ▶ Can you think of a workload where FCFS performs poorly?

FCFS (aka FIFO)

Good scenario

- **Example:** P_1 runs for 100 seconds, while P_2 and P_3 run for 10 seconds

- ▶ **Order of arrival:** $P_2 \rightarrow P_3 \rightarrow P_1$
(all arrived at $t = 0$, with a very small delay among them)
- ▶ **Average turnaround time:** $\frac{10 + 20 + 120}{3} = 50 \text{ seconds}$




FCFS (aka FIFO)

Bad scenario

- **Example:** P_1 runs for 100 seconds, while P_2 and P_3 run for 10 seconds

- ▶ **Order of arrival:** $P_1 \rightarrow P_2 \rightarrow P_3$
(all arrived at $t = 0$, with a very small delay among them)
- ▶ **Average turnaround time:** $\frac{100 + 110 + 120}{3} = 110 \text{ seconds}$



FCFS (aka FIFO)


Convoy effect

- **Convoy effect:** short consumers get queued behind a heavyweight one
 - ▶ Imagine a supermarket queue, and you got stuck behind a customer with 3 full shopping carts...
- **Question:** Can you think of a better policy that reduces turnaround time?

SJF

Shortest Job First


- **Selection criteria:** Run the process with the shortest runtime, then the next shortest, ...
- **Example:** P_1 executes for 100 seconds, while P_2 and P_3 execute for 10 seconds
 - **Order of arrival:** $P_1 \rightarrow P_3 \rightarrow P_2$ (all arrive at $t = 0$, with a very slight delay among them)
Considering this example with SJF, the order of process arrival is irrelevant, right?
 - **Average turnaround time:** $\frac{10 + 20 + 120}{3} = 50 \text{ seconds}$ (~2x reduction over FCFS!)



SJF

Processes arriving at different times

- Lets relax Assumption 2 - *all processes arrive in the system at the same time*
- **Example:** P_1 executes for 100 seconds, while P_2 and P_3 execute for 10 seconds
 - **Order of arrival:** P_1 arrives at instant $t = 0$, while P_2 and P_3 arrive at instant $t = 10$
 - **Average turnaround time:** $\frac{100 + (110 - 10) + (120 - 10)}{3} = 103.33 \text{ seconds}$ (**convoy effect!**)




Preemption

- By definition SJF is a **non-preemptive** scheduler,
 - ▶ This means that processes run until they yield the CPU
 - ▶ To alleviate the previous convoy effect, the scheduler must **preempt P_1** and run instead another process
- Combine the **timer interrupt** and **context switching** mechanisms!
 - ▶ Check the *Scheduling Mechanisms slides!*
- Let's relax Assumption 3 - *Once started, processes runs until completion*

STCF (aka SRT)

Shortest Time-to-Completion First (aka Shortest Remaining Time)

- **Selection criteria:** Anytime a new process enters the system, schedule the process with the least (shortest) runtime left
- **Example:** Process P_1 executes for 100 seconds, while P_2 and P_3 execute for 10 seconds
 - ▶ **Order of arrival:** P_1 arrives at instant $t = 0$, while P_2 and P_3 arrive at instant $t = 10$
 - ▶ STCF must preempt P_1 when P_2 and P_3 arrive
 - ▶ **Average turnaround time:**
$$\frac{(120 - 0) + (20 - 10) + (30 - 10)}{3} = 50 \text{ seconds}$$




STCF (aka SRT)

Another example

- Assuming the arrival and execution times at the table below

- Average turnaround time: $\frac{(16 - 0) + (7 - 2) + (5 - 4) + (11 - 5)}{4} = 7.75 \text{ seconds}$

| Process | Arrival | Execution Time (CPU Burst) |
|---------|---------|-------------------------------|
| 1 | 0 | 7 |
| 2 | 2 | 4 |
| 3 | 4 | 1 |
| 4 | 5 | 4 |



Interactive Systems

Response Time

- With our current assumptions and metric, STCF is a good solution!
 - ▶ Useful for **batch** systems!
- With **time-shared** computers, users started asking for interactive performance
 - ▶ And, thus, a new metric arises: **Response Time**
(i.e., time from when the process arrives in the system to the first time it is scheduled)


$$T_{\text{response}} = T_{\text{firstrun}} - T_{\text{arrival}}$$

tempo que entra no sistema e dá uma resposta

Interactive Systems

Example with SJF


- **Example:** Processes P_1, P_2 and P_3 , each taking 5 seconds to execute
 - ▶ **Order of arrival:** $P_1 \rightarrow P_2 \rightarrow P_3$ (all arrive at $t = 0$, with a very small delay)
 - ▶ **Average response time:** $\frac{0 + 5 + 10}{3} = 5 \text{ seconds}$



RR

Round Robin (aka time-slicing)

- **Selection criteria:** Run a process for a given time slice (**scheduling quantum**), then switch to the next process in queue (queued processes are served in a FCFS fashion)
- **Example:** Processes P_1, P_2 and P_3 , each taking 5 seconds to execute
 - ▶ **Order of arrival:** $P_1 \rightarrow P_2 \rightarrow P_3$ (all arrived at instant 0, with a very small delay)
 - ▶ **Scheduling quantum:** 1 second
 - ▶ **Average response time:** $\frac{0 + 1 + 2}{3} = 1 \text{ second}$



RR


Considerations

- The **scheduling quantum** is critical for RR's efficiency
 - ▶ Short values - better response time
 - ▶ Large values - similar to a FCFS policy
- However, a very small scheduling quantum can be problematic
 - ▶ The cost (time spent) of context switching can dominate the overall performance (more time spent on switching processes than executing processes)
 - ▶ The duration of the scheduling quantum must **amortize** the cost of context switching
- Examples (the values are not supposed to be representative of today's systems):
 - ▶ *scheduling quantum = 10 ms and context-switching = 1 ms* - 10% wasted time
 - ▶ *scheduling quantum = 100 ms and context-switching = 1 ms* - 1% wasted time

RR

Turnaround time

- **Example:** Processes P_1, P_2 and P_3 , each taking 5 seconds to execute
 - ▶ **Order of arrival:** $P_1 \rightarrow P_2 \rightarrow P_3$ (all arrived at $t = 0$, with a very small delay)
 - ▶ **Scheduling quantum:** 1 second
 - ▶ **Average turnaround time:** $\frac{13 + 14 + 15}{3} = 14 \text{ seconds}$




RR

Another example

- Assuming the following arrival and execution times

- Average response time: $\frac{(0) + (2 - 2) + (5 - 4) + (7 - 6)}{4} = 0.5 \text{ seconds}$
- Average turnaround time: $\frac{(16 - 0) + (12 - 2) + (6 - 4) + (15 - 5)}{4} = 9.5 \text{ seconds}$

| Process | Arrival | Execution Time (CPU Burst) |
|---------|---------|-------------------------------|
| 1 | 0 | 7 |
| 2 | 2 | 4 |
| 3 | 4 | 1 |
| 4 | 5 | 4 |



Summary

FCFS, SJF, STCF, RR

● FCFS (FIFO)

- ▶ Simple and easy to implement!
- ▶ Turnaround time is very sensitive to the processes order of arrival

● SJF and STCF

- ▶ Both improve turnaround time. STCF is better when processes arrive at distinct times
- ▶ Both require knowing runtime in advance, and are bad when considering response time

● RR

- ▶ A fair scheduler, good for improving response time (important for interactive systems)
- ▶ One of the worst algorithms for turnaround time (fair algorithms are bad for this metric)

● Tradeoff between optimizing turnaround time and response time

Accounting for I/O


Scheduler decisions

- Lets relax Assumption 4 - processes *only use the CPU (i.e., no I/O is done)*
- Scheduler decisions
 - ▶ If a process blocks for I/O (trap) should another process be scheduled?
 - ▶ When I/O is done (interrupt operation) should the corresponding process be resumed immediately?

Accounting for I/O

Naive approach - without overlap


- **Example:** Processes P_1 and P_2 execute for 50 ms
 - ▶ P_1 executes I/O operations every 10ms, P_2 only uses CPU
- The scheduler runs P_1 and then P_2



Accounting for I/O

Resource overlap

- Let's assume an algorithm based on STCF, but
 - Every 10 ms time slice from P_1 is treated as an independent process
 - Interactive processes (blocking for I/O) run frequently
 - CPU-intensive processes run while interactive ones block for I/O (overlap)
 - Better turnaround and response times (**Homework:** calculate them!)



MLFQ (Multi-Level Feedback Queue)

No more Oracle

- Lets relax our last *Main Assumption - the runtime of processes is known*
- OS goals
 - ▶ Optimize turnaround time, like in SJF or STCF
(but without knowing the runtime of processes...)
 - ▶ Make the system feel responsive to interactive users, like in RR
(good response time)
- Solution
 - ▶ The scheduler must combine multiple scheduling policies
 - ▶ The scheduler must learn from the past to make better decisions in the future!


Note: This work led to the award of a Turing Award to its author Fernando J. Corbató

MLFQ

Basic rules

- Multiple queues, ordered according to their **priority**
- Priorities are used to choose what process runs next
 - ▶ **Rule 1:** If Priority(P_1) > Priority(P_2), P_1 runs (P_2 does not)
 - ▶ **Rule 2:** If Priority(P_1) = Priority(P_3), P_1 & P_3 run in RR

quando tenho mais que um processo na mesma fila aplica-se o RR




- How are priorities assigned to processes and, do these change over time?
 - ▶ Compensate well behaved processes that relinquish CPU frequently
 - ▶ Demote processes that are CPU intensive
 - ▶ **Allotment:** the amount of time a process can spend at a given priority level
 - The time slice for each queue (priority level) may change (e.g., longer time slices for low priority queues)

MLFQ

Queue movement


- **Rule 3:** When a process enters the system it has the highest priority (top queue)
- **Rule 4:** Once a process uses up its time allotment at a given level, its priority is reduced (i.e., moves down one queue)



MLFQ

Preemption of low priority processes


- When new processes, with higher priority arrive, low priority processes executing are preempted
- P_1 is preempted (does not run) until P_2 has the same priority (both are at Q_1)



MLFQ

Gaming


- Interactive processes (e.g., P_2) may relinquish the CPU (block for I/O) before their allotment expires. Should their priority be kept?
 - ▶ Bad users could trick (*game*) the OS and starve other processes by writing program that repeatedly use 99% of the allotment and then yield the CPU



MLFQ

Better Accounting


- **Updated Rule 4:** Once a process uses up its time allotment at a given level (**regardless of how many times it has given up the CPU**), its priority is reduced (i.e., moves down one queue)
 - ▶ Thus, the allotment time must consider the sum of all CPU bursts of a process



MLFQ

Starvation


- Starvation of processes with longer runtimes may still exist



MLFQ

Priority Boost (aging mechanism)

- Rule 5: After a time period S, move all the jobs to the topmost queue



MLFQ

Summary

- Short duration processes run with high priority (approximates STCF)
 - ▶ Good for turnaround time
- High priority processes are frequently switched (approximates RR)
 - ▶ Good for response time
- Challenge: several tuning knobs in MLFQ
 - How many queues?
 - How long should the allotment time per queue be?
 - When should the priority boost be called?
 - ▶ No easy answer
 - One must know the OS and workloads running there

More Information

- **Chapters 7 and 8** - Remzi H. Arpaci-Dusseau, Andrea C. Arpaci-Dusseau. **Operating Systems: Three Easy Pieces**. Arpaci-Dusseau Books, 2018.
- Avi Silberschatz, Peter Baer Galvin, Greg Gagne. **Operating System Concepts (10. ed)**. John Wiley & Sons, 2018.

Proportional-Sharing

Interesting topics - not covered in detail in this course

- Proportional-share (aka fair-share) schedulers

- ▶ **Goal:** give each job a fair percentage of CPU (**fairness metric**)
- ▶ Interesting algorithms
 - Lottery Scheduling
 - Stride Scheduling
 - Completely Fair Scheduler (CFS)

- **CFS** is the most widely-used fair scheduler

- ▶ Used by default in several Linux distributions
- ▶ Goals
 - Be fair among the multiple processes ready to run
 - Be very efficient on choosing the next process to run

CFS

A very brief overview

- **Virtual runtime (*vruntime*):** counts the CPU time used by a given process
 - ▶ Incremented every time a process uses the CPU
- **Selection criteria:** choose the process with the lowest *vruntime*
- *vruntime* may be incremented differently for specific processes
 - **Example:** time increases slower for high priority processes
 - Priority may be chosen by users and system administrators (**nice** level of a process)
- A **red-black tree** structure is used to index processes by their *vruntime* and to efficiently choose the next to run

CFS

A very brief overview

- **Scheduling latency target (`sched_latency`):** Within the `sched_latency` time interval each process should have a fair chance to use the CPU
 - ▶ **Example:** Assuming 4 processes and `sched_latency = 48 ms`, each process is given a time slice of 12 ms
- The time-slice adjusts dynamically when processes arrive/leave the system
 - ▶ **Let's go back to the previous example:** if 2 processes finish executing, the time slice for each of the two remaining processes is increased to 24 ms
- A minimum time slice must be set when a large number of processes are waiting to be scheduled (**`min_granularity`**)
 - ▶ **Question:** Do you know why?

Advanced topics

Not covered in this course

● Multiprocessor scheduling

- ▶ Modern computers can have several CPUs...
- ▶ New challenges:
 - Synchronization
 - Cache Affinity
 - ...

● Want to know more about these topics?

- **Chapters 9 and 10 - Remzi H. Arpaci-Dusseau, Andrea C. Arpaci-Dusseau. Operating Systems: Three Easy Pieces.** Arpaci-Dusseau Books, 2018.

Questions?